



Internship Project Report

DermalScan: AI Facial Skin Aging Detection App

Submitted By:

Kamsali Niharika

Under guidance of Mentor **Mr. Praveen**

Infosys Springboard Virtual Internship

Problem Statement:

Develop an AI-powered facial skin–aging detection system using EfficientNetB0 that classifies aging indicators such as **wrinkles**, **dark spots**, **puffy eyes**, and **clear skin** from uploaded images.

The pipeline includes:

- Face detection via Haar Cascades
- Preprocessing & augmentation
- Deep learning classification
- Web-based visualization with bounding boxes
- Prediction export and logging
-

Objectives:

- Build a deep learning classifier using **EfficientNetB0** with at least **90% accuracy**.
- Detect aging regions and output **percentage-based predictions**.
- Create a **Streamlit UI** for real-time inference (≤ 5 seconds).
- Prepare dataset → preprocess → augment → train → evaluate → deploy.

Milestone – 1: Dataset Preparation & Preprocessing

Module – 1: Dataset Setup and Image Labeling

The dataset was manually curated and organized into four classes:

- **puffy_eyes**
- **wrinkles**
- **dark_spots**
- **clear_skin**

Each image was placed into the corresponding folder and renamed in a structured format (e.g., puffy_eyes_1.jpg, wrinkles_42.jpg) using an automated Python renaming script.

1. Image Counting per Class

Sample Code:

```
for cls in CLASSES:
```

```
    folder = DATA_DIR / cls  
    count = len(list(folder.glob("*.*")))  
    print(cls, ":", count)
```

Purpose: ensure that the dataset is correctly loaded and classes are balanced.

2. Class Distribution Visualization

Sample Code:

```
sns.barplot(x=df_counts.index, y=df_counts['count'])  
plt.title("Number of Images per Class")  
plt.xlabel("Class")  
plt.ylabel("Image Count")  
plt.show()
```

Output:




Fig – 1

3. Sample Image Visualization

Sample Code:

```
show_samples("clear skin", n=6)
```

Output:



This helped confirm:


- Images were placed in correct categories

4. Image Dimension Analysis

Sample Code:

```
sns.histplot(df_sizes["width"], kde=True, label="Width")  
sns.histplot(df_sizes["height"], kde=True, label="Height")  
plt.legend()  
plt.title("Image Width & Height Distribution")  
plt.show()
```

Output:



Ensured consistency in image shapes before resizing.

5. Brightness Distribution per Class

A KDE plot was generated to analyze lighting differences between classes:

Code:

```
sns.kdeplot(data=df_bright, x="brightness", hue="class", fill=True)
```

Output:




Fig – 4

Module – 2: Image Preprocessing and Augmentation

Images were resized, normalized, augmented, and label – encoded as required by EfficientNetB0.

1. Resizing & Normalizing (224×224):

Sample Code:

```
img = Image.open(img_path).convert("RGB")
img = img.resize((224, 224))
img = np.array(img) / 255.0
```

Output:

X shape: (908, 224, 224, 3)

y shape: (908, 4)

2. One-Hot Encoding of Labels

Sample Code:

```
y_encoded = tf.keras.utils.to_categorical(labels, num_classes=4)
```

3. Data Augmentation

Sample Code:

```
datagen = ImageDataGenerator(  
    rotation_range=15,  
    zoom_range=0.1,  
    horizontal_flip=True  
)  
  
aug_iter = datagen.flow(sample_img, batch_size=1)  
plt.imshow(next(aug_iter)[0])
```

Output:

Augmentation Examples



Fig – 5 (Rotation, Zoom, Flip)

6. Augmentation Quality Visualization

Sample Code:

```
plt.figure(figsize=(8,4))

sns.histplot(X.ravel(), bins=50, color="blue", label="Original", stat="density")

aug_batch = datagen.flow(X, y, batch_size=100)
augmented_sample, _ = next(aug_batch)

sns.histplot(augmented_sample.ravel(), bins=50, color="red", label="Augmented",
stat="density")

plt.legend()
plt.title("Pixel Intensity Distribution Before vs After Augmentation")
plt.show()
```

Output:




Fig – 6

After preprocessing and augmentation, the final dataset was split into training, validation, and testing subsets using the stratified train-test split method to maintain class balance. An 80/10/10 ratio was used, which is widely accepted for deep learning tasks. All six final arrays (`X_train`, `y_train`, `X_val`, `y_val`, `X_test`, `y_test`) were saved for efficient loading during model training.

Milestone – 2: Model Training & Evaluation

Module – 3: Model Development & Training

1. Objective:

To build a reliable deep-learning model that classifies:
Wrinkles, Dark Spots, Puffy Eyes, Clear Skin

2. Dataset Improvement:

To ensure richer learning and better generalization:

- Increased dataset from ~300 → **~500 images per class**
- Final usable dataset ≈ **1800+ images**
- Balanced all classes
- Removed noisy, tiny & corrupt images
- Standardized input to **224×224 resolution**

3. Model Selection:

Chosen Model: EfficientNet (Fine-Tuned)

Reason:

- Excellent accuracy – efficiency balance
- Strong feature learning for skin textures
- Stable convergence

4. Training Strategy:

Phase – 1:

Freeze EfficientNet → Train Classification Head

Phase – 2:

Unfreeze selected layers → Fine-tune at lower LR

5. Training Configuration:

- Input Size: 224×224
- Optimizer: Adam
- Batch Size: 32
- Epochs: 50+
- Loss: Categorical Crossentropy
- Regularization: Dropout + EarlyStopping


6. Performance:

- Training Accuracy: $\approx 99\%$
- Validation Accuracy: $\approx 94\%$
- Stable curves (no heavy overfitting)
- Strong confusion matrix behavior

Final Model Selected → EfficientNet Fine-Tuned

7. Model Comparison Table:


Model Variant	Train Acc	Val Acc	Epochs	Batch
EfficientNet – Phase 1	~93%	~90%	20	32
EfficientNet – Fine Tune	~99%	~94%	40	32





Classification Report:

Label	Precision	Recall	F1-score	Support
clear skin	0.98	0.95	0.96	100
dark spots	0.91	0.95	0.93	96
puffy eyes	0.95	0.93	0.94	80
wrinkles	0.93	0.94	0.93	83
accuracy			0.94	359
macro avg	0.94	0.94	0.94	359
weighted avg	0.94	0.94	0.94	359



Module – 4: Facial Region Detection & Prediction Pipeline

1. Objective:

To develop a robust computer vision pipeline that automatically detects human faces within an image, extracts the Region of Interest (ROI), and passes it through a multi-stage deep learning inference engine for skin condition classification and age estimation.

2. Face Detection Algorithm

- **Technique:** Haar Feature-based Cascade Classifier.
- **Model:** haarcascade_frontalface_default.xml
- **Implementation Details:**
 - **Scale Factor:** 1.2 (Compensates for faces appearing smaller/larger due to distance).
 - **Min Neighbors:** 6 (High threshold to eliminate false positives like fabric patterns or shadows).
 - **Aspect Ratio Filter:** Implemented a geometric filter $0.7 < (w/h) < 1.3$ to strictly reject non-face rectangles.

4. Age Estimation:

- Implemented **real integer age prediction**.
- Returns values like **21, 22, 23...**

5. Final AI Prediction Output:

For each face system outputs:

- Wrinkles %
- Dark Spots %
- Puffy Eyes %
- Clear Skin %
- **Dominant Condition Highlighted**
- Predicted Integer Age

Displayed on image with clean black bounding box + label.

AGE : 28
WRINKLES : 2.46%
PUFFY EYES: 95.79%
DARK SPOTS: 13.61%
CLEAR SKIN: 71.26%



Milestone – 3: System Integration & Prototype

Module – 5: Frontend Development (Streamlit UI)

1. Objective:

To develop a responsive, user-friendly web interface that allows users to upload images, visualize real-time predictions, and supports batch processing.

2. UI Design & Layout:

- **Framework:** Streamlit (Python).
- **Theme:** "Neon/Cyberpunk" (Dark Mode) for high-contrast visibility of skin conditions.
- **Layout:** Wide layout with a sidebar for configuration and a main area for batch image grids.
- **Interactive Elements:**
 - Drag-and-drop file uploader (Supports: JPG, PNG, JPEG).
 - Real-time processing status indicators.
 - Expandable sections (st.expander) for detailed probability breakdowns.

3. Frontend Implementation:

The UI was built to handle multiple images simultaneously. The layout dynamically adjusts columns based on the number of uploaded files.

Sample Code Snippet (UI Layout):

```
st.set_page_config(page_title="AI DermalScan Pro", layout="wide")
st.markdown("""
<style>
.stApp { background-color: #050505; color: #ffffff; }
h1 { color: #00ffcc; text-shadow: 0 0 10px #00ffcc; }
</style>
""", unsafe_allow_html=True)
```

4. Dashboard Visualization: The interface provides immediate visual feedback. Users can see the original image overlaid with analysis data side-by-side with statistical charts.




Fig: AI DermalScan Pro Landing Page with File Uploader.

Module – 6: Backend Integration & Smart Logic

1. Objective:

To integrate the trained **MobileNetV2** (Skin Model) and **AgeNet** (Caffe Model) into a unified inference pipeline that processes images in real-time.

2. Face Detection & Preprocessing:

- **Algorithm:** Haar Cascade Classifier (Strict Mode).
- **Optimization:** Implemented scaleFactor=1.2 and minNeighbors=5 to reduce false positives (e.g., detecting clothes as faces).
- **Context Padding (Innovation):** Standard face detection crops too tightly, causing age prediction errors (predicting adults as babies). We implemented a **20% Context Padding** logic to include the forehead and chin for accurate age estimation.

Sample Code Snippet (Context Padding):

```
# Zoom out 20% to capture hairline and chin for Age Model  
pad_w = int(w * 0.20)  
pad_h = int(h * 0.20)  
face_roi_bgr_padded = img_bgr[y1_pad:y2_pad, x1_pad:x2_pad]
```

3. Smart Bio-Age Algorithm: A custom logic layer was added to correct the "Age 5 Error" common in Caffe models.

- **Rule 1 (Imperfection Ban):** If Wrinkles or Puffy Eyes are detected with >40% confidence, "Child" age buckets (0-12) are mathematically suppressed.
- **Rule 2 (Selfie Detection):** If the face occupies >30% of the image width, the minimum age floor is raised to 18.
- **Rule 3 (Severity Penalty):** Adds +2 to +7 years to the predicted age based on the confidence of the wrinkle detection.




Fig: Accurate Age & Wrinkle Detection using Smart Bio-Age Logic.

4. 3D Visualization & Reporting:

Instead of flat bar charts, we integrated **Plotly 3D Donut Charts** to visualize the probability distribution of skin conditions. The dominant condition is "exploded" (pulled out) for emphasis.

Sample Code Snippet (3D Chart):

```
fig = go.Figure(data=[go.Pie(  
    labels=labels, values=values, pull=pull_values, hole=0.4,  
    marker=dict(colors=NEON_COLORS)  
)])
```




Fig: 3D Donut Chart visualizing probability of Puffy Eyes vs Clear Skin.

5. Final Output & Export:

The system generates a fully annotated image with a drop-shadow text overlay for readability. Users can download individual processed images or a batch CSV report.

Filename	Face ID	Condition	Age	Confidence	Clear Skin %	Dark Spots %	Puffy Eyes %	Wrinkles %
shutterstock_1072	1	Wrinkles	85	99.80%	0.00%	0.20%	0.00%	99.80%

Fig: Batch Analysis Report and CSV Export Feature.

Conclusion:

The prototype successfully integrates the Frontend and Backend. The **Context Padding** and **Smart Bio-Age** algorithms significantly improved age prediction accuracy compared to the raw Caffe model outputs. The system is now ready for final deployment.